
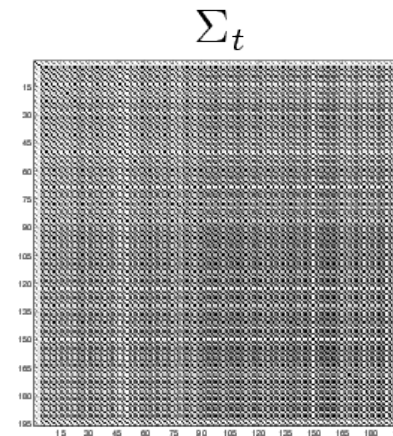


# L15. POSE-GRAPH SLAM





NA568 Mobile Robotics: Methods & Algorithms

# Today's Topic



- Nonlinear Least Squares
- Pose-Graph SLAM
- Incremental Smoothing and Mapping



# Feature-Based SLAM




Filtering Problem:  
Motion Prediction Causes Fill-in



Conceptual Abstraction: Full-SLAM



Marginalize out the Poses



$$\begin{bmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{bmatrix} = \begin{bmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{bmatrix} - \begin{bmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{bmatrix} \Lambda_{LL} - \begin{bmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{bmatrix} \Lambda_{Lx} - \begin{bmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{bmatrix} \Lambda_{xx} - \begin{bmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{bmatrix} \Lambda_{xL}$$



# Feature-Based SLAM

- In feature-based SLAM, the information matrix fills in unless we enforce **approximations** to make it sparse
  - ▣ This is because we are continually marginalizing out the robot trajectory from the state representation
- What if we were to marginalize out the landmarks instead?




# Pose-Graph SLAM

- Feature-based SLAM requires approximations to enforce sparsity.
- Furthermore, this approximation is non-trivial.



Conceptual Abstraction: Full-SLAM

Marginalize out the Landmarks



$$\begin{bmatrix} \text{Sparse Matrix} \end{bmatrix} = \begin{bmatrix} \Lambda_{xx} \end{bmatrix} - \begin{bmatrix} \Lambda_{xL} \end{bmatrix} \begin{bmatrix} \Lambda_{LL} \end{bmatrix}^{-1} \begin{bmatrix} \Lambda_{Lx} \end{bmatrix}$$

Key idea: marginalizing out the landmarks preserves locality

# Three Main SLAM Paradigms

Kalman  
filter

Particle  
filter

Graph-  
based



**least squares  
approach to SLAM**


# Least Squares in General

- Approach for computing a solution for an **overdetermined system**
- “More equations than unknowns”
- Minimizes the **sum of the squared errors** in the equations
- Standard approach to a large set of problems

**Today: Application to SLAM**

# Graph-Based SLAM


- Constraints connect the poses of the robot while it is moving
- Constraints are inherently uncertain





# Graph-Based SLAM

- Observing previously seen areas generates constraints between non-successive poses




# Idea of Graph-Based SLAM

- Use a **graph** to represent the problem
- Every **node** in the graph corresponds to a pose of the robot during mapping
- Every **edge** between two nodes corresponds to a spatial constraint between them
- **Graph-Based SLAM:** Build the graph and find a node configuration that minimizes the error introduced by the constraints

# Graph-Based SLAM in a Nutshell


- Every node in the graph corresponds to a robot position and a laser measurement
- An edge between two nodes represents a spatial constraint between the nodes



KUKA Halle 22, courtesy of P. Pfaff

# Graph-Based SLAM in a Nutshell


- Every node in the graph corresponds to a robot position and a laser measurement
- An edge between two nodes represents a spatial constraint between the nodes



KUKA Halle 22, courtesy of P. Pfaff

# Graph-Based SLAM in a Nutshell

- Once we have the graph, we determine the most likely map by correcting the nodes




# Graph-Based SLAM in a Nutshell

- Once we have the graph, we determine the most likely map by correcting the nodes  
... like this



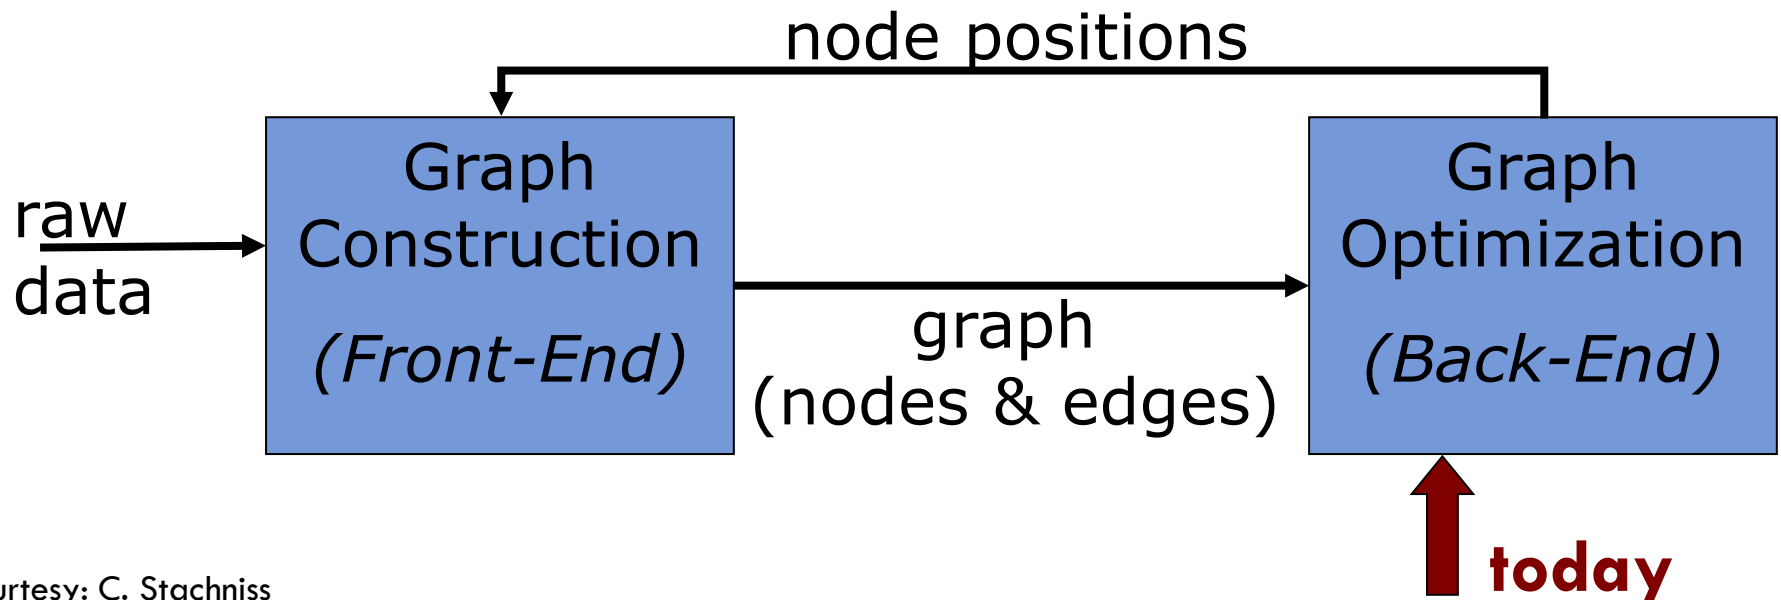
# Graph-Based SLAM in a Nutshell

- Once we have the graph, we determine the most likely map by correcting the nodes  
... like this
- Then, we can render a map based on the known poses



# The Overall SLAM System


- Interplay of front-end and back-end
- Map helps to determine constraints by reducing the search space
- Topic today: optimization






# The Graph

- It consists of  $n$  nodes  $\mathbf{X} = \mathbf{X}_{1:n}$
- Each  $\mathbf{X}_i$  is a 2D or 3D transformation (the pose of the robot at time  $t_i$ )
- A constraint/edge exists between the nodes  $\mathbf{X}_i$  and  $\mathbf{X}_j$  if...



# Create an Edge If... (1)


- ...the robot moves from  $\mathbf{x}_i$  to  $\mathbf{x}_{i+1}$
- Edge corresponds to odometry




The edge represents the **odometry** measurement

# Create an Edge If... (2)

- ...the robot observes the same part of the environment from  $\mathbf{x}_i$  and from  $\mathbf{x}_j$




Measurement from  $\mathbf{x}_i$



Measurement from  $\mathbf{x}_j$

# Create an Edge If... (2)

- ...the robot observes the same part of the environment from  $\mathbf{x}_i$  and from  $\mathbf{x}_j$
- Construct a **virtual measurement** about the position of  $\mathbf{x}_j$  seen from  $\mathbf{x}_i$



Edge represents the position of  $\mathbf{x}_j$  seen from  $\mathbf{x}_i$  based on the **observation**

# Transformations

- Transformations can be expressed using **homogenous coordinates**
- Odometry-Based edge


$$(\mathbf{X}_i^{-1} \mathbf{X}_{i+1})$$

- Observation-Based edge

$$(\mathbf{X}_i^{-1} \mathbf{X}_j)$$

How node  $i$  sees node  $j$

# Simultaneous Localization and Mapping



Given a **single camera** feed,  
estimate the 3D **position of the camera** and  
the 3D **positions of all landmark** points in the world

# Visual SLAM: Why Filter?

Image and Vision Computing 30 (2012) 65–77



Contents lists available at [SciVerse ScienceDirect](#)

Image and Vision Computing

journal homepage: [www.elsevier.com/locate/imavis](http://www.elsevier.com/locate/imavis)



Editors Choice Article

## Visual SLAM: Why filter? ☆

Hauke Strasdat <sup>a,\*</sup>, J.M.M. Montiel <sup>b</sup>, Andrew J. Davison <sup>a</sup>

<sup>a</sup> Department of Computing, Imperial College London, UK

<sup>b</sup> Instituto de Investigacion en Ingenieria de Aragon (I3A), Universidad de Zaragoza, Spain

### ARTICLE INFO

#### Article history:

Received 2 August 2011

Received in revised form 13 December 2011

Accepted 17 February 2012

#### Keywords:

SLAM

Structure from motion

Bundle adjustment

EKF

Information filter

Monocular vision


Stereo vision

### ABSTRACT

While the most accurate solution to off-line structure from motion (SfM) problems is undoubtedly to extract as much correspondence information as possible and perform batch optimisation, sequential methods suitable for live video streams must approximate this to fit within fixed computational bounds. Two quite different approaches to real-time SfM – also called visual SLAM (simultaneous localisation and mapping) – have proven successful, but they sparsify the problem in different ways. Filtering methods marginalise out past poses and summarise the information gained over time with a probability distribution. Keyframe methods retain the optimisation approach of global bundle adjustment, but computationally must select only a small number of past frames to process. In this paper we perform a rigorous analysis of the relative advantages of filtering and sparse bundle adjustment for sequential visual SLAM. In a series of Monte Carlo experiments we investigate the accuracy and cost of visual SLAM. We measure accuracy in terms of entropy reduction as well as root mean square error (RMSE), and analyse the efficiency of bundle adjustment versus filtering using combined cost/accuracy measures. In our analysis, we consider both SLAM using a stereo rig and monocular SLAM as well as various different scenes and motion patterns. For all these scenarios, we conclude that keyframe bundle adjustment outperforms filtering, since it gives the most accuracy per unit of computing time.


© 2012 Elsevier B.V. All rights reserved.

# Visual SLAM





# The SLAM Problem ( $t=0$ )




## Onboard sensors:

- Wheel odometry
- Inertial measurement unit (gyro, accelerometer)
- Sonar
- Laser range finder
- Camera
- RGB-D sensors


# The SLAM Problem ( $t=1$ )

Odometry measurement

Robot



Landmark measurement



Landmark 1




Landmark 2

# The SLAM Problem ( $t=n-1$ )

Odometry measurement


Robot

Landmark  
measurement




# The SLAM Problem ( $t=n$ )

Odometry measurement



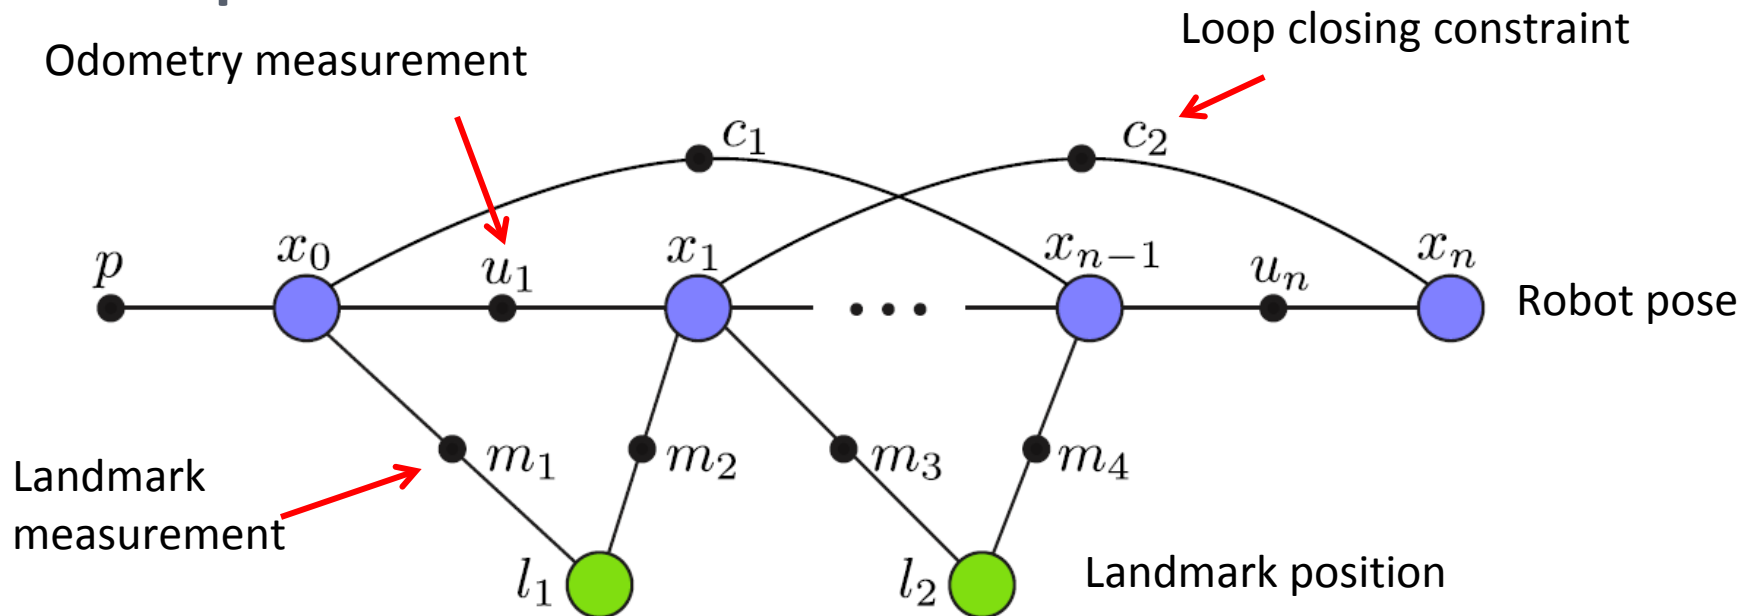
# Factor Graph Representation




Bipartite graph with ***variable nodes*** and ***factor nodes***




# Factor Graph Representation: Pose Graph




Bipartite graph with ***variable nodes*** and ***factor nodes***




# Factor Graph Representation: Bundle Adjust.



Bipartite graph with ***variable nodes*** and ***factor nodes***



# Nonlinear Least-Squares




Gaussian noise

$$\operatorname{argmin}_{\mathbf{x}} \sum_i \|h_i(\mathbf{x})\|_{\mathbf{E}}^2$$

Repeatedly solve linearized system (GN)

$$\operatorname{argmin}_{\mathbf{x}} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2$$


$$\mathbf{A} = \begin{bmatrix} F_{11} & & G_{11} & & \\ F_{12} & & & G_{12} & \\ F_{13} & & & & G_{13} \\ & F_{21} & G_{21} & & \\ & F_{22} & & G_{22} & \\ & F_{23} & & & G_{23} \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} \xi_1 \\ \xi_2 \\ \delta_1 \\ \delta_2 \\ \delta_3 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} b_{11} \\ b_{12} \\ b_{13} \\ b_{14} \\ b_{15} \\ b_{16} \end{bmatrix}$$






# Solving the Linear Least-Squares System

Solve:  $\operatorname{argmin}_x \|Ax - b\|^2$



Normal equations

$$A^T A x = A^T b$$



Information matrix


Measurement Jacobian

# Solving the Linear Least-Squares System

- Can we simply invert  $A^T A$  to solve for  $x$ ?


Normal equations

$$A^T A x = A^T b$$



Information matrix


- Yes, but we shouldn't...  
The inverse of  $A^T A$  is dense  $\rightarrow O(n^3)$
- Can do much better by taking advantage of sparsity!



# Solving the Linear Least-Squares System

[Dellaert and Kaess, IJRR 06]


Solve:  $\operatorname{argmin}_x \|Ax - b\|^2$



Measurement Jacobian

Normal equations


$$A^T A x = A^T b$$



Information matrix

Matrix factorization


$$A^T A = R^T R$$




Square root information matrix

# Matrix – Square Root Factorization

- QR on  $A$ : Numerically More Stable




- Cholesky on  $A^T A$ : Faster




# Retaining Sparsity: Variable Ordering

Fill-in depends on elimination order:




factor  
→




Default ordering (poses, landmarks)

↓ permute



factor  
→




Ordering based on COLAMD heuristic [Davis04]  
(best order: NP hard)

# Solving by Backsubstitution

After factorization:  $R^T R \mathbf{x} = A^T \mathbf{b}$


- Forward substitution

$R^T \mathbf{y} = A^T \mathbf{b}$ , solve for  $\mathbf{y}$




- Backsubstitution

$R \mathbf{x} = \mathbf{y}$ , solve for  $\mathbf{x}$



# Full Bundle Adjustment




From Strasdat et al, 2011 IVC “Visual SLAM: Why filter?”

- Graph grows with time:
  - ▣ Have to solve a sequence of increasingly larger BA problems
  - ▣ Will become too expensive even for sparse Cholesky

F. Dellaert and M. Kaess, “Square Root SAM: Simultaneous localization and mapping via square root information smoothing,” IJRR 2006


# Keyframe Bundle Adjustment



- Drop subset of poses to reduce density/complexity
- Only retain “keyframes” necessary for good map
- Complexity still grows with time, just slower




# Filter



- Keyframe idea not applicable: map would fall apart
- Instead, marginalize out previous poses
  - ▣ Extended Kalman Filter (EKF)
- Problems when used for Visual SLAM:
  - ▣ All points become fully connected → expensive
  - ▣ Relinearization not possible → inconsistent

# Incremental Solver

- Back to full BA and keyframes:



- New information is added to the graph
- Older information does not change
- Can be exploited to obtain an efficient solution!

# Incremental Smoothing and Mapping (iSAM)

# iSAM

Solving a growing system:

- ▣ Exact/batch (quickly gets expensive)
- ▣ Approximations
- ▣ Incremental Smoothing and Mapping (iSAM)


New measurements ->

Key idea:


- ▣ Append to existing matrix factorization
- ▣ “Repair” using Givens rotations

Periodic batch steps for

- ▣ Relinearization
- ▣ Variable reordering (to keep sparsity)



# Factor Updates with Givens Rotations



Old  $R$  factor

New rows

New  $R$  factor:  
Triangulated using  
Givens Rotations

Constant time!

Zeroing entries with Givens  
Rotations:


$$\begin{array}{c} \text{row } k \\ \text{row } i \end{array} \begin{array}{|c|} \hline \begin{array}{cccc} \cdot & 1 & & \\ & c & & s \\ & & 1 & \cdot \\ -s & & & 1 & c \end{array} \\ \hline \text{Givens} \end{array} \cdot \begin{array}{|c|} \hline \begin{array}{c} \text{gray triangle} \\ x \end{array} \\ \hline R \end{array} = \begin{array}{|c|} \hline \begin{array}{c} \text{gray triangle} \\ \text{red bar} \\ 0 \end{array} \\ \hline R' \end{array}$$

**Numerically stable !**

# Variable Reordering – Constrained COLAMD

## Greedy approach

Arbitrary placement of newest variable



Number of affected variables:


low

high



## Constrained Ordering


Newest variables forced to the end



**Much cheaper!**

# Variable Reordering – Fill-in


Incremental ordering still yields good overall ordering



- Only slightly more fill-in than batch COLAMD ordering
- Constrained ordering is worse than naïve/greedy:
  - Suboptimal ordering because of partial constraint, but cheaper to update!

# iSAM

Example from real sequence:  
Square root inf. matrix →  
Side length: 21000 variables  
Dense: 1.7GB, sparse: 1MB





# Next Lecture

- FastSLAM